

# COP4610 Project 1

A Simple Shell

# Overview

- A simple shell that interprets commands from either user input or a file
  - Interactive vs Batch modes
- Commands are delimited by ‘;’
- All valid commands on a given line of input should be run simultaneously
  - `fork()` and `wait()`
- Shells exit when either given the built-in command “quit” or when EOF is reached
  - However, the shell must wait for other commands to finish running

# Interactive Mode

- Interactive mode is the default shell behavior; i.e., when the program is run without any arguments
  - `$ tinysh`
- Interactive mode displays a prompt to the user whenever the shell is ready for more input
  - `prompt> command1; command2; ...`
- When EOF is reached (CTRL+D in this mode) or the “quit” command is entered
  - First, ensure completion of all commands on the same line in which “quit” appears
  - Once all commands have finished execution, exit

# Batch Mode

- Batch mode is entered only when the shell is invoked with an argument
  - `$ tinysh /path/to/valid/batch/file`
- Do not forget that “quit” is a valid command in batch mode
  - When “quit” is encountered in batch mode, do **not** handle any subsequent lines
  - Commands that appear on the same line as “quit” must still be executed
- Batch files will be 0 or more lines of command strings, i.e.
  - `$ cat example.batch`  
`ls -l; echo hello`  
`echo goodbye; quit; cat /foo/bar`  
`cat /bar/foo` <= NOT EXECUTED

# Library Routines: fgets()

- <http://linux.die.net/man/3/fgets>
- `#include <stdio.h>`
- `char *fgets(char *s, int num, FILE *stream);`
- Reads characters from stream until
  - num - 1 characters are read
  - Newline is encountered
  - EOF is reached
- Returns str on success
- Returns NULL
  - On error
  - When EOF reached and nothing read

- Example usage

```
int BUFSIZE = 1024;
char buf[BUFSIZE];

if(fgets(buf, BUFSIZE, stdin) == NULL) {
    // error handling goes here
} else {
    printf("buf: %s\n", buf);
}
```

# Library Routines: strtok()

- <http://linux.die.net/man/3/strtok>
- `#include <string.h>`
- `char strtok(char *str, const char *delim);`
- `delim` acts as the delimiter for tokens
  - That is, the symbol that will be matched to split a string into tokens
- `str` should be the string to tokenize on the **first** call to `strtok()`
  - Subsequent calls wishing to tokenize the same string will then pass `NULL` in place of `str`

- Example usage

```
char* token;  
char* delim = "+";
```

```
token = strtok("foo+ bar+ foobar", delim);  
while(token != NULL) {  
    // do something with token  
    token = strtok(NULL, delim);  
}
```

- In the above code, `token` should take the value of
  - `foo`, `bar`, `foobar`

# Library Routines: fork()

- <http://linux.die.net/man/2/fork>
- `#include <unistd.h>`
- `pid_t fork(void);`
- Duplicates calling process
  - Functionally, this *child* is an exact duplicate of the *parent*
  - Unique PID
- To the parent, `fork()` returns either -1 on failure or the PID of the child on success
  - 0 is returned to the child
- Typically, the parent process will `wait()` for a child to finish execution

- Example usage

```
pid_t pid;
if((pid = fork()) < 0) {
    perror("fork");
} else if(pid == 0) {
    // inside the child, so do work and..
    exit(1);
}
```

# Library Routines: wait()

- <http://linux.die.net/man/2/wait>
- `#include <sys/types.h>`
- `#include <sys/wait.h>`
- **`pid_t wait(int *status);`**
- Waits for a state change in a child of the calling process
  - Termination or start/stop due to a signal
- Allows for child's resources to be released
  - Without wait()ing, child becomes a zombie
- This function blocks until the child's status has changed
  - Returns immediately if the child has already terminated

- Example usage

// in a parent process after a fork()

```
pid_t pid;
```

```
int status;
```

```
pid = wait(&status);
```

```
printf("PID %d exited with status %d\n",  
pid, status);
```



# Library Routines: execvp()

- <http://linux.die.net/man/3/execvp>
- `#include <unistd.h>`
- **`int execvp(const char *file, char *const argv[]);`**
- file is the name of a file to be executed
  - I.e., a command
- argv[] is the full array of arguments
  - argv[0] holds the filename, and can be used as the file argument
  - argv[] must be terminated by a null pointer
- On success, execvp() does not return
  - The execvp()'d command replaces the current process image
  - Therefore, an error with execvp() has occurred if this function is returned from

- Example usage

```
char* args[3];  
args[0] = "ls";  
args[1] = "-l";  
args[2] = NULL;
```

```
if(/* successful fork */) {  
    execvp(args[0], args);  
    printf("execvp() failed\n");  
    exit(0);  
}
```

# Hints

- Depending on your implementation, you may consider keeping track of how many processes were fork()ed
  - You will need to wait() on each of these processes **after** they have all been fork()ed
- Test your program as you go
  - If you write a couple lines of code and think of something that might go wrong
    - Try it out and see if your program breaks
    - Include error handling
      - Waiting to include error handling may cause hours of headache-inducing debugging